ABHINAV REJA

+1 (480) 942-8069 | Tempe, AZ | abhinavreja2000@gmail.com | Portfolio | LinkedIn | GitHub

SUMMARY

Software Engineering graduate with 2+ years of experience in full-stack systems, APIs, and LLM applications. Built Next.js products, PostgreSQL services, and tested AI workflows from design through deployment.

EDUCATION

M.S. in Software Engineering — **May 2026**
Arizona State University, Arizona, USA — **GPA: 4.0/4.0**

Bachelor's in Information Technology — **May 2022**
Oriental Institute of Science and Technology, Bhopal, India — **GPA: 3.36/4.0**

EXPERIENCE

Forward Deployed Engineer, **OZZI (OZO Development LLC)**, Phoenix, AZ — **June 2026 - Present**

- **Travel Insurance:** Shipped a Vercel-deployed Next.js/TypeScript chatbot with Claude and Supabase PostgreSQL for policy Q&A and flight-delay claims, verifying delays via AviationStack and applying deterministic payout rules; reduced LLM cost **71%** using Anthropic prompt caching.
- Mapped stateless claim state from conversation history for each request, with deterministic business logic and repeatable testing across distributed server instances.
- Validated AI chatbot releases with 42 claim regression cases, 95 conversational scenarios, and PII-redaction gates to catch response quality, claim-logic, and privacy failures before release.
- **Campaign Sign Planning:** Engineered a recommendation engine that ranks high-visibility campaign-sign locations across ~11,000 ADOT traffic points using PostGIS, OpenStreetMap road modeling, spacing rules, and district-anchored geocoding.

Backend Developer Intern, **Vosyn Inc.**, Chicago, IL — **June 2025 - August 2025**

- Developed a history-aware content recommendation API powered by a fine-tuned Llama-3-8B (LoRA) classifier across 50+ genres, using batch inference and cached lookups to **improve CTR by 15% (A/B tested)**.
- Redesigned the PostgreSQL schema for audio/video/text content tables using composite indexes and normalized relations, **reducing p90 retrieval latency by ~30%** across 3+ downstream services.
- Improved bulk audio and video ingestion with a Django REST batch-processing API, **reducing end-to-end ingestion time by 11%** by decoupling heavy media pipelines from synchronous request workflows.
- Implemented a Java/Cucumber integration testing framework in the GitHub Actions CI/CD pipeline for API and end-to-end workflow validation, **reducing regression testing time per release cycle by 15%**.

Software Engineer, **WINspect Technologies Pvt. Ltd.**, Mumbai, India — **January 2023 - May 2024**

- Built C++ REST APIs for a time-sync dashboard serving 500+ distributed clients at **sub-200ms p95 latency**, then added Redis TTL caching that **cut average read latency 50%** under high-frequency polling.
- Automated 40+ integration tests for synchronization features using Google Test and wired them into the Jenkins CI/CD pipeline, **reducing manual regression testing effort per release cycle by 30%**.
- Wrote Python automation to troubleshoot and remediate recurring production failures, **cutting ticket resolution time by 18%**.

PROJECTS

WIAL Platform | Opportunity Hack 2026 | Demo • Code — **March 2026**

- Created a user-facing Next.js app with cross-lingual OpenAI embeddings and pgvector search matching coaches across **20+ chapters in 10+ languages**; unified fragmented WordPress sites with **sub-4s loads on 3G**.

CrodeX | HackASU 2025 Winner, Claude Developer Tools Track | Devpost — **November 2025**

- Launched an LLM-powered developer tool that migrates Flask, Django, and Spring Boot projects while preserving application logic, selected over **100+ submissions**.
- Designed a multi-stage agentic pipeline using Gemini and Claude APIs to analyze backend project structure, extract dependency graphs, and generate framework migrations **in under 90 seconds per project**, with human-in-the-loop review flagging ambiguous dependency conflicts before code generation.

Face Recognition on AWS | IaaS, Serverless, and Edge — **September 2025 - November 2025**

- Provisioned an IaaS pipeline with a custom autoscaler, S3/SQS queueing, and PyTorch workers on a custom AMI, achieving **0.96s latency and 80% recognition accuracy** across 100 concurrent requests.
- Delivered a serverless variant (ECR, MTCNN, SQS-triggered FaceNet inference, **1.78s latency**) and an edge variant deployed on AWS Greengrass with MQTT messaging and Lambda-backed recognition (**0.78s latency**).

TECHNICAL SKILLS

Languages: TypeScript, JavaScript, Python, SQL, Java, C++
Backend & Frameworks: Next.js, React, Node.js, Django, Flask, FastAPI, Spring Boot
Cloud/DevOps: AWS (EC2, S3, SQS, Lambda, ECR), Google Cloud Run, Docker, Jenkins, GitHub Actions, Vercel
Databases: PostgreSQL, Supabase, pgvector, PostGIS, MongoDB, Redis
AI/ML: Claude API, OpenAI API, Gemini API, Llama-3 (LoRA) inference, RAG, prompt caching, agentic pipelines, PyTorch
Testing: Google Test, Cucumber, regression suites as deploy gates